


Pairwise alignment of DNA/protein sequences

Yuzhen Ye (yye@indiana.edu)
School of Informatics & Computing, IUB

We compare biological molecules, not any two strings!

- Sequence alignment reveals function, structure, and evolutionary information!
- From edit distance to distance between two biological molecules with biological meaning—the scoring matrix
- Local alignment versus global alignment
- But still the dynamic programming algorithm is the algorithm behind pairwise alignment of biological sequences

Aligning DNA Sequences: scoring matrix



Simple scoring

- When mismatches are penalized by $-\mu$, indels are penalized by $-\sigma$, and matches are rewarded with $+1$, the resulting score is:

$$\#matches - \mu(\#mismatches) - \sigma(\#indels)$$

Scoring matrices

To generalize scoring, consider a $(4+1) \times (4+1)$ **scoring matrix** δ .

In the case of an amino acid sequence alignment, the scoring matrix would be a $(20+1) \times (20+1)$ size. The addition of 1 is to include the score for comparison of a gap character "-".

This will simplify the algorithm as follows:

$$s_{ij} = \max \begin{cases} s_{i-1,j-1} + \delta(v_i, w_j) \\ s_{i-1,j} + \delta(v_i, -) \\ s_{i,j-1} + \delta(-, w_j) \end{cases}$$

Scoring matrices for DNA sequence alignment

- A simple positive score for matches and a negative for mismatches and gaps are most often used.
- Transversions penalized more than transitions
 - transitions: replacement of a purine base with another purine or replacement of a pyrimidine with another pyrimidine (A \leftrightarrow G, C \leftrightarrow T)
 - transversions: replacement of a purine with a pyrimidine or vice versa.
 - Transition mutations are more common than transversions

Making a scoring matrix for protein sequence alignment

- Scoring matrices are created based on biological evidence.
- Alignments can be thought of as two sequences that differ due to mutations.
- Some of these mutations have little effect on the protein's function, therefore some penalties, $\delta(v_i, w_j)$, will be less harsh than others.
- We need to know how often one amino acid is substituted for another in related proteins

Scoring matrix: example

	A	R	N	K
A	5	-2	-1	-1
R	-	7	-1	3
N	-	-	7	0
K	-	-	-	6

AKRANR

KAAANK

$$-1 + (-1) + (-2) + 5 + 7 + 3 = 11$$

- Notice that although R and K are different amino acids, they have a positive score.
- Why? They are both positively charged amino acids → will not greatly change function of protein.

Conservation

- Amino acid changes that tend to preserve the physico-chemical properties of the original residue
 - Polar to polar
 - aspartate → glutamate
 - Nonpolar to nonpolar
 - alanine → valine
 - Similarly behaving residues
 - leucine to isoleucine

Common scoring matrices for protein sequence alignment

- Amino acid substitution matrices
 - PAM
 - BLOSUM
- Try to compare protein coding regions at amino acid level
 - DNA is less conserved than protein sequences (codon degeneracy; synonymous mutations)
 - Less effective to compare coding regions at nucleotide level

Reading: Chapter 4, 4.3

PAM

- Point Accepted Mutation (Dayhoff et al.)
- 1 PAM = PAM1 = 1% average change of all amino acid positions
 - After 100 PAMs of evolution, not every residue will have changed
 - some residues may have mutated several times
 - some residues may have returned to their original state
 - some residues may not have changed at all

PAM1 & PAM250

Substitution	PAM1	PAM250
Phe to Ala	0.0002	0.04
Phe to Arg	0.0001	0.01
Phe to Asp		
...		
Phe to Phe	0.9946	0.32
...		
Sum	1	1

Normalized probability scores for changing Phe to other amino acids at PAM1 and PAM250 evolutionary distances

Chapter 3, table 3.2

Log-odds substitution matrices

- Using amino acid changes that were observed in closely related proteins; they represented amino acid substitutions that don't significantly change the structure and function of the protein.
 - "accepted mutations" or "accepted" by natural selection.
- Log-odds of the probability of matching a pair of amino acids in this database relative to a random one
 - Ref: Amino acid substitution matrices from protein blocks (PNAS. 1992, 89(22): 10915–10919)

Log odd scores

Define f_{ij} as the frequency of observing amino acid pair i, j . Then the observed probability of occurrence for each i, j pair is

$$q_{ij} = f_{ij} / \sum_{i=1}^{20} \sum_{j=1}^{20} f_{ij}$$

The probability of occurrence of the i th amino acid in an i, j pair is,

$$p_i = q_{ii} + \sum_{j \neq i} q_{ij} / 2,$$

The *expected* probability of occurrence e_{ij} for each i, j pair is computed as,

$$e_{ij} = p_i p_j \quad \text{if } i = j \\ = 2p_i p_j \quad \text{else}$$

A lod (logarithm of odd) is then calculated in bit units as, $s_{ij} = \log_2(q_{ij}/e_{ij})$

The amino acid substitution is considered as a Markov model

- A Markov model is characterized by a series of changes of state in a system such that a change from one state to another does not depend on the previous history of the state
- Use of the Markov model makes it possible to extrapolate amino acid substitutions observed over a relatively short period of evolutionary time to longer periods of evolutionary time
 - $\text{PAM}_x = \text{PAM}_1^x$
 - The multiplication of two PAM1 matrices \rightarrow PAM2

PAM250

- PAM250 is a widely used scoring matrix:
- $\text{PAM250} = \text{PAM1}^{250}$

	Ala	Arg	Asn	Asp	Cys	Gln	Glu	Gly	His	Ile	Leu	Lys	...
A	R	N	D	C	Q	E	G	H	I	L	K	...	
Ala A	13	6	9	9	5	8	9	12	6	8	6	7	...
Arg R	3	17	4	3	2	5	3	2	6	3	2	9	...
Asn N	4	4	6	7	2	5	6	4	6	3	2	5	...
Asp D	5	4	8	11	1	7	10	5	6	3	2	5	...
Cys C	2	1	1	1	52	1	1	2	2	2	1	1	...
Gln Q	3	5	5	6	1	10	7	3	7	2	3	5	...
...													
Trp W	0	2	0	0	0	0	0	0	1	0	1	0	...
Tyr Y	1	1	2	1	3	1	1	1	3	2	2	1	...
Val V	7	4	4	4	4	4	4	4	5	4	15	10	...

BLOSUM

- Blocks Substitution Matrix
- Scores derived from observations of the frequencies of substitutions in **blocks of local alignments in related proteins**
- Matrix name indicates evolutionary distance
 - BLOSUM62 was created using sequences sharing no more than 62% identity

BLOSUM versus PAM

- The PAM family
 - PAM matrices are based on **global** alignments of closely related proteins.
 - The PAM1 is the matrix calculated from comparisons of sequences with no more than 1% divergence; Other PAM matrices are **extrapolated** from PAM1.
- The BLOSUM family
 - BLOSUM matrices are based on **local** alignments (blocks)
 - All BLOSUM matrices are based on **observed** alignments (BLOSUM 62 is a matrix calculated from comparisons of sequences with no less than 62% divergence)
- Higher numbers in the PAM matrix naming scheme denote larger evolutionary distance; BLOSUM is the opposite.
 - For alignment of distant proteins, you use PAM150 instead of PAM100, or BLOSUM50 instead of BLOSUM62.

Local vs. global alignment

- Global Alignment

```
--T--CC-C-AGT--TATGT-CAGGGGACACG-A-GCATGCAGA-GAC
      ||| ||| ||| ||| ||| ||| ||| ||| ||| ||| ||| |||
AATTGCCGCC-GTCGT-T-TTCAG----CA-GTTATG-T-CAGAT--C
```

- Local Alignment—better alignment to find conserved segment

```
          tccCAGTTATGTCAGgggacacgagcatgcagagac
          |||||
aattgccgccgtcgttttcagCAGTTATGTCAGatc
```

Local vs. global alignment

- The Global Alignment Problem tries to find the longest path between vertices $(0,0)$ and (n,m) in the edit/alignment graph.
 - The Needleman–Wunsch algorithm
- The Local Alignment Problem tries to find the longest path among paths between **arbitrary vertices** (i,j) and (i',j') in the edit graph.
 - The Smith–Waterman algorithm


Local alignments: why?

- Two genes in different species may be similar over short conserved regions and dissimilar over remaining regions.
- Example:
 - Homeobox genes have a short region called the homeodomain that is highly conserved between species.
 - A global alignment would not find the homeodomain because it would try to align the ENTIRE sequence

The local alignment problem

- Goal: Find the best local alignment between two strings
- Input: Strings \mathbf{v} , \mathbf{w} and scoring matrix δ
- Output: Alignment of substrings of \mathbf{v} and \mathbf{w} whose alignment score is maximum among all possible alignment of all possible substrings


Local alignment: example



Compute a "mini" Global Alignment to get Local

Local alignment: running time

- Long run time $O(n^6)$:



- In the grid of size $n \times n$ there are $\sim n^2$ vertices (i,j) that may serve as a source and $\sim n^2$ vertices (i',j') that may serve as a sink.
- For each such vertices computing alignments from (i,j) to (i',j') takes $O(n^2)$ time.

We do NOT go with this algorithm!

The local alignment recurrence

- The largest value of s_{ij} over the whole edit graph is the score of the best local alignment.

- The recurrence:

$$s_{ij} = \max \begin{cases} 0 \\ s_{i-1,j-1} + \delta(v_i, w_j) \\ s_{i-1,j} + \delta(v_i, -) \\ s_{i,j-1} + \delta(-, w_j) \end{cases}$$

Notice there is only this change from the original recurrence of a Global Alignment

The local alignment recurrence

- The largest value of s_{ij} over the whole edit graph is the score of the best local alignment.
- The recurrence:

$$s_{ij} = \max \begin{cases} 0 \\ s_{i-1,j-1} + \delta(v_i, w_j) \\ s_{i-1,j} + \delta(v_i, -) \\ s_{i,j-1} + \delta(-, w_j) \end{cases}$$

Power of ZERO: there is only this change from the original recurrence of a Global Alignment - since there is only one "free ride" edge entering into every vertex


- Complexity: $O(N^2)$, or $O(MN)$

→ **Initialization will be different**

Scoring indels: naive approach

- A fixed penalty σ is given to every indel:
 - $-\sigma$ for 1 indel,
 - -2σ for 2 consecutive indels
 - -3σ for 3 consecutive indels, etc.
- Can be too severe penalty for a series of 100 consecutive indels

Arbitrary gap penalty?




There are many such edges!

Adding them to the graph increases the running time of the alignment algorithm by a factor of n (where n is the number of vertices)

So the complexity increases from $O(n^2)$ to $O(n^3)$

Affine gap penalties

- In nature, a series of k indels often come as a single event rather than a series of k single nucleotide events:



Affine gap penalties

- Score for a gap of length x is:
 $-(\rho + \sigma x)$
 where $\rho > 0$ is the penalty for introducing a gap:
 gap opening penalty
 ρ will be large relative to σ :
 gap extension penalty
 because you do not want to add too much of a penalty for extending the gap.
- Reduced penalties (as compared to naïve scoring) are given to runs of horizontal and vertical edges

Affine gap penalty recurrences

$$\text{gap}(L) = \sigma + \rho * L$$

$$D(i, j) = \max \begin{cases} D(i-1, j) + \sigma & \text{Continue gap in } y \text{ (deletion)} \\ S(i-1, j) + (\sigma + \rho) & \text{Start gap in } y \text{ (deletion)} \end{cases}$$

$$I(i, j) = \max \begin{cases} I(i, j-1) + \sigma & \text{Continue gap in } x \text{ (insertion)} \\ S(i, j-1) + (\sigma + \rho) & \text{Start gap in } x \text{ (insertion)} \end{cases}$$

$$S(i, j) = \max \begin{cases} S(i-1, j-1) + \delta(x_i, y_j) & \text{Match or mismatch} \\ I(i, j) & \text{End with deletion} \\ D(i, j) & \text{End with insertion} \end{cases}$$

Readings

- Chapter 4, 4.1
 - “Alignment can reveal homology between sequences” (similarity vs homology)
 - “It is easier to detect homology when comparing protein sequences than when comparing nucleic acid sequences”
- Primer article: [What is dynamic programming](#) by Eddy

We will continue on

- Significance of an alignment (score)
 - Homologous or not?
- Faster alignment tools for database search